

Handwritten Recognition via Convolutional Neural Network

May 30, 2019

Jian Sun
DUID: 873397832

1 Content

- Introduction
- Model
- Result
- Train VGG-16 with MNIST from Scratch
- Predict MNIST by Trained Imagenet Weight in VGG-16
- Freeze All The Convolutional Layers Retrain VGG-16 Network On MNIST
- Freeze All The Fully Connected Layers Retrain VGG-16 Network On MNIST
- Discussion
- Appendix

2 Introduction

In this project, we implement two methods to recognize MNIST dataset.

One is to code the VGG-16 from scratch and train it with the MNIST from the very beginning, another method is to train MNISR with existing VGG-16 weight drawn from Imagenet dataset. Specifically, we will discuss the prediction capability after freezing all convolutional layers or freezing all fully connected layers. This will help us to learn the power of different layers in convolutional neural network. Before doing the experiment, we hypothesize that fully connected layer is vital crucial. Without it, the whole nerual network will perform extremely worse. Conversely, convolional layer doesn't play an important layer for MNIST dataset, since it is a binery image dataset, not the colorful one.

3 Model

3.1 Data Processing

The whole MNIST is splitted as the following three parts.

The training set has 48000 28X28 images;
the validation set has 12000 28X28 images;
the testing set has 10000 28X28 images.

Then, given that the VGG-16 doesn't accept image which size is smaller than 32X32, the current images have to be enlarged. The method here is adding paddings. And the final size is 34X34, which fits for VGG-16 model.

3.2 Train VGG-16 with MNIST from Scratch

This part starts from stacking layers. With the help of `Sequence()` from keras, we build the network up. And we set the batch size as 16, epoch as 500, optimizer as adam, loss as categorical cross entropy loss.

Callbacks function is used here to reduce the learning rate automatically and stop the training when validation loss doesn't change for 10 epochs, this function saves us lots of time and is very convenient.

3.3 Predict MNIST by Trained Imagenet Weight in VGG-16

A very powerful function, `VGG16()` from `keras.applications` helps to construct the neural network here. And we set the batch size as 16, epoch as 100, optimizer as adam, loss as categorical cross entropy loss.

Callbacks function is also used here to reduce the learning rate automatically and stop the training when validation loss doesn't change for 8 epochs, this function saves us lots of time and is very convenient.

3.4 Freeze All The Convolutional Layers Retrain VGG-16 Network On MNIST

A very powerful function, `VGG16()` from `keras.applications` helps to construct the neural network here. And we set the batch size as 16, epoch as 100, optimizer as adam, loss as categorical cross entropy loss. Additionally, the convolutional layers are frozen by code: `layer.trainable = False`.

Callbacks function is also used here to reduce the learning rate automatically and stop the training when validation loss doesn't change for 8 epochs, this function saves us lots of time and is very convenient.

3.5 Freeze All The Fully Connected Layers Retrain VGG-16 Network On MNIST

A very powerful function, `VGG16()` from `keras.applications` helps to construct the neural network here. And we set the batch size as 16, epoch as 100, optimizer as adam, loss as categorical cross entropy loss. Additionally, the fully connected layers are frozen by code: `layer.trainable = False`.

Callbacks function is also used here to reduce the learning rate automatically and stop the training when validation loss doesn't change for 8 epochs, this function saves us lots of time and is very convenient.

4 Train VGG-16 with MNIST from Scratch

The training stops at epoch 11 with the training loss of 4.9966e-04, the training accuracy of 0.9999, the validation loss of 0.0688 and the validation accuracy of 0.9934, when the learning rate is 7.999999979801942e-07.

Next, we use this weight to predict testing set, the prediction loss is 0.0492 and the prediction accuracy is 0.9953.

```

In [111]: print('Got %d / %d correct' % (num_correct, num_test))
          print('Accuracy = %f' % (np.mean(y_test == y_pred)))
          print(confusion_matrix(y_test, y_pred))
          print(classification_report(y_test, y_pred,
                                     target_names=list(label_dict.values()),digits=3))

          plt.figure(figsize=(8,8))
          cnf_matrix = confusion_matrix(y_test, y_pred)
          classes = list(label_dict.values())
          plt.imshow(cnf_matrix, interpolation='nearest')
          plt.colorbar()
          tick_marks = np.arange(len(classes))
          _ = plt.xticks(tick_marks, classes, rotation=90)
          _ = plt.yticks(tick_marks, classes)

```

Got 9953 / 10000 correct


Accuracy = 0.995300

```

[[ 977    0    0    0    0    0    2    1    0    0]
 [   0 1130    1    2    0    0    0    1    1    0]
 [   0    0 1029    1    0    0    0    2    0    0]
 [   0    0    0 1006    0    3    0    0    1    0]
 [   0    0    0    0  977    0    1    0    0    4]
 [   0    0    0    5    0  886    1    0    0    0]
 [   4    2    0    0    0    0  952    0    0    0]
 [   0    2    1    0    0    0    0 1024    0    1]
 [   0    0    0    1    0    0    0    0  972    1]
 [   0    0    0    0    4    1    0    2    2 1000]]

```

	precision	recall	f1-score	support
0	0.996	0.997	0.996	980
1	0.996	0.996	0.996	1135
2	0.998	0.997	0.998	1032
3	0.991	0.996	0.994	1010
4	0.996	0.995	0.995	982
5	0.996	0.993	0.994	892
6	0.996	0.994	0.995	958
7	0.994	0.996	0.995	1028
8	0.996	0.998	0.997	974
9	0.994	0.991	0.993	1009
accuracy				0.995 10000
macro avg				0.995 10000
weighted avg				0.995 10000



5 Predict MNIST by Trained Imagenet Weight in VGG-16

The training stops at epoch 37 with the training loss of 0.0303, the training accuracy of 0.9920, the validation loss of 0.0724 and the validation accuracy of 0.9778, when the learning rate is $6.400000529538374e-08$.

Next, we use this weight to predict testing set, the prediction loss is 0.0672 and the prediction accuracy is 0.9788.

```
In [0]: print('Got %d / %d correct' % (num_correct, num_test))
        print('Accuracy = %f' % (np.mean(y_test == y_pred)))
        print(confusion_matrix(y_test, y_pred))
        print(classification_report(y_test, y_pred,
                                     target_names=list(label_dict.values()), digits=3))
```

```

plt.figure(figsize=(8,8))
cnf_matrix = confusion_matrix(y_test, y_pred)
classes = list(label_dict.values())
plt.imshow(cnf_matrix, interpolation='nearest')
plt.colorbar()
tick_marks = np.arange(len(classes))
_ = plt.xticks(tick_marks, classes, rotation=90)
_ = plt.yticks(tick_marks, classes)

```

Got 9788 / 10000 correct

Accuracy = 0.978800

```


[[ 969   0   1   0   0   4   5   1   0   0]
 [   0 1123   0   0   3   0   4   4   1   0]
 [   1   1 1008   4   2   4   4   5   2   1]
 [   0   0  10  977   0  16   0   3   2   2]
 [   0   1   1   0  967   1   2   3   4   3]
 [   2   0   7   9   1  864   0   2   6   1]
 [   7   0   4   0   0   2  944   0   1   0]
 [   0   5   4   2   8   0   0 1004   1   4]
 [   0   0   6   2   3   4   1   1  956   1]
 [   3   0   3   2   7   2   0   5  11  976]]

```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.987	0.989	0.988	980
1	0.994	0.989	0.992	1135
2	0.966	0.977	0.971	1032
3	0.981	0.967	0.974	1010
4	0.976	0.985	0.980	982
5	0.963	0.969	0.966	892
6	0.983	0.985	0.984	958
7	0.977	0.977	0.977	1028
8	0.972	0.982	0.977	974
9	0.988	0.967	0.977	1009

accuracy			0.979	10000
macro avg	0.979	0.979	0.979	10000
weighted avg	0.979	0.979	0.979	10000



6 Freeze All The Convolutional Layers Retrain VGG-16 Network On MNIST

The training stops at epoch 22 with the training loss of 0.0272, the training accuracy of 0.9924, the validation loss of 0.0723 and the validation accuracy of 0.9788, when the learning rate is $6.400000529538374 \times 10^{-8}$.

Next, we use this weight to predict testing set, the prediction loss is 0.0693 and the prediction accuracy is 0.9771.

```
In [0]: print('Got %d / %d correct' % (num_correct, num_test))
        print('Accuracy = %f' % (np.mean(y_test == y_pred)))
        print(confusion_matrix(y_test, y_pred))
```

```

print(classification_report(y_test, y_pred,
                           target_names=list(label_dict.values()),digits=3))

plt.figure(figsize=(8,8))
cnf_matrix = confusion_matrix(y_test, y_pred)
classes = list(label_dict.values())
plt.imshow(cnf_matrix, interpolation='nearest')
plt.colorbar()
tick_marks = np.arange(len(classes))
_ = plt.xticks(tick_marks, classes, rotation=90)
_ = plt.yticks(tick_marks, classes)

```

Got 9771 / 10000 correct

Accuracy = 0.977100

```


[[ 967    0    1    0    0    3    5    1    2    1]
 [   0 1125    0    0    3    0    4    2    1    0]
 [   1    2 1000    8    2    6    3    6    3    1]
 [   0    0    9  976    0   18    0    4    2    1]
 [   0    1    0    0  968    1    3    2    3    4]
 [   2    0    4   13    1  861    3    2    5    1]
 [   7    1    2    0    1    3  941    0    2    1]
 [   0    3    4    3   10    0    0 1002    1    5]
 [   0    0    3    4    2    4    2    2  955    2]
 [   3    0    3    3    5    3    0    4   12  976]]

```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.987	0.987	0.987	980
1	0.994	0.991	0.993	1135
2	0.975	0.969	0.972	1032
3	0.969	0.966	0.968	1010
4	0.976	0.986	0.981	982
5	0.958	0.965	0.961	892
6	0.979	0.982	0.981	958
7	0.978	0.975	0.976	1028
8	0.969	0.980	0.974	974
9	0.984	0.967	0.976	1009

accuracy			0.977	10000
macro avg	0.977	0.977	0.977	10000
weighted avg	0.977	0.977	0.977	10000



7 Freeze All The Fully Connected Layers Retrain VGG-16 Network On MNIST

The training stops at epoch 9 with the training loss of 2.4916, the training accuracy of 0.1330, the validation loss of 2.5009 and the validation accuracy of 0.1352, when the learning rate is $8.000000525498762e-06$.

Next, we use this weight to predict testing set, the prediction loss is 2.4918 and the prediction accuracy is 0.1296.

```
In [0]: print('Got %d / %d correct' % (num_correct, num_test))
        print('Accuracy = %f' % (np.mean(y_test == y_pred)))
        print(confusion_matrix(y_test, y_pred))
```



```

print(classification_report(y_test, y_pred,
                           target_names=list(label_dict.values()),digits=3))

plt.figure(figsize=(8,8))
cnf_matrix = confusion_matrix(y_test, y_pred)
classes = list(label_dict.values())
plt.imshow(cnf_matrix, interpolation='nearest')
plt.colorbar()
tick_marks = np.arange(len(classes))
_ = plt.xticks(tick_marks, classes, rotation=90)
_ = plt.yticks(tick_marks, classes)

```

Got 1296 / 10000 correct

Accuracy = 0.129600

```

[[ 0  0  0  0  0  0  0  0 980  0]
 [ 0  0  0  4 821 266  0  2  34  8]
 [ 0  0  0  0  5  0  0 10 1015  2]
 [ 0  0  0  0  4  0  0  0 1005  1]
 [ 0  0  0  5 315  27  0 12  612 11]
 [ 0  0  0  0  8  0  0  2  881  1]
 [ 0  0  0  0  8  1  0  3  939  7]
 [ 0  0  0  0 217  48  0  8  702 53]
 [ 0  0  0  0  2  0  0  0  972  0]
 [ 0  0  0  0 15  5  0  0  988 1]]

```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------


0	0.000	0.000	0.000	980
1	0.000	0.000	0.000	1135
2	0.000	0.000	0.000	1032
3	0.000	0.000	0.000	1010
4	0.226	0.321	0.265	982
5	0.000	0.000	0.000	892
6	0.000	0.000	0.000	958
7	0.216	0.008	0.015	1028
8	0.120	0.998	0.214	974
9	0.012	0.001	0.002	1009

accuracy			0.130	10000
macro avg	0.057	0.133	0.050	10000
weighted avg	0.057	0.130	0.049	10000

```

/usr/local/lib/python3.6/dist-packages/sklearn/metrics/classification.py:1437: UndefinedMetricWarning: Precision is undefined because no predicted samples were equal to any of the class labels
'precision', 'predicted', average, warn_for)

```



7.1 Discussion

Here is a detailed table for all 4 parts.

Methods	Scratch	Retrain	Freeze Conv	Freeze Fully Con
Train Acc	99.99%	99.20%	99.24%	13.30%
Train Loss	5e ⁻⁴	0.0303	0.0272	2.4916
Valid Acc	99.34%	97.78%	97.88%	13.52%
Valid Loss	0.0688	0.0724	0.0723	2.5009
Test Acc	99.53%	97.88%	97.71%	12.96%
Test Loss	0.0492	0.0672	0.0693	2.4918
Time Per Epoch	143s	5s	5.5s	3s
Learning Rate	8e ⁻⁷	6e ⁻⁸	6e ⁻⁸	8e ⁻⁶

As we can see, the total rank for training part is Scratch > Freeze Conv > Retrain > Freeze Fully Con. Scratch performs best on training part. VGG-16 freezing all convolutional layers performs better than Retrain it on MNIST. It is a little weird, but it is truth. Meanwhile, VGG-16 freezing all full connected layers performs worst.

The total rank for testing part is Scratch > Retrain > Freeze Conv > Freeze Fully Con.

In general, from this comparison, we recognize that full connected layer is so important for CNN, however, if we close convolutional layer, our result will not be influenced a lot. Perhaps, it is due to MNIST dataset as the binary image dataset, not the colorful one. Thus losing convolutional layers doesn't affect a lot on CNN's performance. However, losing fully connected layer means different, it will ruin your prediction, since we need softmax to do classification.

CNN is a powerful model, there is huge CNN families. We hope to dig more.

8 Appendix

8.1 Load Packages

```
In [0]: # import dataset and separate them as train set and test set
        # index x represents image, index y represents label
```

```
import tensorflow as tf
import keras
from numpy.linalg import *
from tensorflow import keras
import matplotlib.pyplot as plt
from keras.optimizers import Adam
from keras.models import Sequential
from keras.applications import VGG16
from keras.utils import to_categorical
from keras.layers.convolutional import *
from keras import callbacks, layers, optimizers, models
import os, cv2, random, sklearn, sklearn.metrics, numpy as np
from sklearn.preprocessing import StandardScaler
from keras.metrics import categorical_crossentropy
from sklearn.model_selection import train_test_split
from keras.layers.normalization import BatchNormalization
from keras.layers import Input, Dense, Conv2D, MaxPooling2D, Activation, Dropout, Flatten
from sklearn.metrics import classification_report, confusion_matrix
from keras.layers.convolutional import Convolution2D, MaxPooling2D, ZeroPadding2D
from __future__ import absolute_import, division, print_function, unicode_literals
```


8.2 Load Dataset and Data Preparing

```
In [0]: # download MNIST dataset from keras
        (x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
        y_train_labels = to_categorical(y_train)
        y_test_labels = to_categorical(y_test)
        # make sure the 10 classes
```

```
label_dict = {0: '0', 1: '1', 2: '2', 3: '3', 4: '4',
              5: '5', 6: '6', 7: '7', 8: '8', 9: '9'}
```

In [86]: *# Verify that the data is in the correct format.*

```
plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(x_train[i], cmap=plt.cm.binary)
    plt.xlabel(label_dict[y_train[i]])
plt.show()
```



```

In [186]: x_train = x_train.reshape(np.shape(x_train)[0], 28*28)
          x_test = x_test.reshape(np.shape(x_test)[0], 28*28)
          # Convert the images into 3 channels
          X_train=np.dstack([x_train] * 3)
          X_test=np.dstack([x_test] * 3)
          print(np.shape(X_train))
          print(np.shape(X_test))
          # Reshape images as per the tensor format required by tensorflow
          X_train = X_train.reshape(-1, 28,28,3)
          X_test = X_test.reshape (-1,28,28,3)
          print(np.shape(X_train))
          print(np.shape(X_test))
          # Resize the images 34*34 as required by VGG16
          X_tr = np.pad(X_train, ((0,0),(3,3),(3,3),(0,0)), 'constant')
          X_te = np.pad(X_test, ((0,0),(3,3),(3,3),(0,0)), 'constant')
          X_tr = np.float32(X_tr)/255.
          X_te = np.float32(X_te)/255.
          print(np.shape(X_tr))
          print(np.shape(X_te))

```

```

(60000, 784, 3)
(10000, 784, 3)
(60000, 28, 28, 3)
(10000, 28, 28, 3)
(60000, 34, 34, 3)
(10000, 34, 34, 3)

```

```

In [187]: # Splitting train data as train and validation data
          train_X,valid_X,train_label,valid_label = train_test_split(X_tr,
                                                                    y_train_labels,
                                                                    test_size=0.2,
                                                                    random_state=13)
          # Finally check the data size whether it is as per tensorflow and VGG16 requirement
          train_X.shape,valid_X.shape,train_label.shape,valid_label.shape

```

```

Out[187]: ((48000, 34, 34, 3), (12000, 34, 34, 3), (48000, 10), (12000, 10))

```

```

In [0]: # Define the parameters for instantiating VGG16 model.
        IMG_WIDTH = 34
        IMG_HEIGHT = 34
        IMG_DEPTH = 3
        BATCH_SIZE = 16

```

8.3 Train VGG-16 with MNIST from Scratch

```

In [176]: input_tensor = Input(shape=(IMG_WIDTH, IMG_HEIGHT, IMG_DEPTH))
          #input_tensor = input_tensor
          vg16_m = VGG16(weights=None,input_tensor = input_tensor)

```

```

type(vg16_m)
model = Sequential()
for layer in vg16_m.layers[0:22]:
    model.add(layer)
model.summary()

```

Layer (type)	Output Shape	Param #
block1_conv1 (Conv2D)	(None, 34, 34, 64)	1792
block1_conv2 (Conv2D)	(None, 34, 34, 64)	36928
block1_pool (MaxPooling2D)	(None, 17, 17, 64)	0
block2_conv1 (Conv2D)	(None, 17, 17, 128)	73856
block2_conv2 (Conv2D)	(None, 17, 17, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
block4_conv1 (Conv2D)	(None, 4, 4, 512)	1180160
block4_conv2 (Conv2D)	(None, 4, 4, 512)	2359808
block4_conv3 (Conv2D)	(None, 4, 4, 512)	2359808
block4_pool (MaxPooling2D)	(None, 2, 2, 512)	0
block5_conv1 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv2 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv3 (Conv2D)	(None, 2, 2, 512)	2359808
block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0
flatten (Flatten)	(None, 512)	0
fc1 (Dense)	(None, 4096)	2101248

```

-----
fc2 (Dense)                (None, 4096)                16781312
=====
Total params: 33,597,248
Trainable params: 33,597,248
Non-trainable params: 0
-----

```

```

In [140]: model.add(Dense(10, activation='softmax'))
          model.summary()

```

```

-----
Layer (type)                Output Shape                Param #
=====
block1_conv1 (Conv2D)       (None, 34, 34, 64)         1792
-----
block1_conv2 (Conv2D)       (None, 34, 34, 64)         36928
-----
block1_pool (MaxPooling2D)  (None, 17, 17, 64)         0
-----
block2_conv1 (Conv2D)       (None, 17, 17, 128)        73856
-----
block2_conv2 (Conv2D)       (None, 17, 17, 128)        147584
-----
block2_pool (MaxPooling2D)  (None, 8, 8, 128)         0
-----
block3_conv1 (Conv2D)       (None, 8, 8, 256)          295168
-----
block3_conv2 (Conv2D)       (None, 8, 8, 256)          590080
-----
block3_conv3 (Conv2D)       (None, 8, 8, 256)          590080
-----
block3_pool (MaxPooling2D)  (None, 4, 4, 256)         0
-----
block4_conv1 (Conv2D)       (None, 4, 4, 512)          1180160
-----
block4_conv2 (Conv2D)       (None, 4, 4, 512)          2359808
-----
block4_conv3 (Conv2D)       (None, 4, 4, 512)          2359808
-----
block4_pool (MaxPooling2D)  (None, 2, 2, 512)         0
-----
block5_conv1 (Conv2D)       (None, 2, 2, 512)          2359808
-----
block5_conv2 (Conv2D)       (None, 2, 2, 512)          2359808
-----
block5_conv3 (Conv2D)       (None, 2, 2, 512)          2359808
-----

```

```

-----
block5_pool (MaxPooling2D)   (None, 1, 1, 512)        0
-----
flatten (Flatten)            (None, 512)               0
-----
fc1 (Dense)                  (None, 4096)             2101248
-----
fc2 (Dense)                  (None, 4096)             16781312
-----
dense_10 (Dense)             (None, 10)               40970
=====
Total params: 33,638,218
Trainable params: 33,638,218
Non-trainable params: 0
-----

```

```
In [0]: NB_EPOCHS = 100
```

```

# Compile the model.
model.compile(Adam(lr=0.0001),
               loss='categorical_crossentropy',
               metrics=['accuracy'])

```

```

In [105]: from keras import callbacks
# Incorporating reduced learning and early stopping for callback
reduce_learning = callbacks.ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.2,
    patience=2,
    verbose=1,
    mode='auto',
    epsilon=0.0001,
    cooldown=2,
    min_lr=0)

early_stopping = callbacks.EarlyStopping(
    monitor='val_loss',
    min_delta=0.0003,
    patience=6,
    verbose=1,
    mode='auto')

callbacks = [reduce_learning, early_stopping]

```

```

/usr/local/lib/python3.6/dist-packages/keras/callbacks.py:1065: UserWarning: `epsilon` argument
warnings.warn("`epsilon` argument is deprecated and '

```



```
In [106]: mt=model.fit(train_X, train_label,
                        batch_size=16, epochs=NB_EPOCHS,
                        verbose=2, callbacks=callbacks,
                        validation_data=(valid_X, valid_label)
                        )
```

Train on 48000 samples, validate on 12000 samples

Epoch 1/100

- 147s - loss: 0.0311 - acc: 0.9933 - val_loss: 0.0585 - val_acc: 0.9894

Epoch 2/100

- 144s - loss: 0.0332 - acc: 0.9933 - val_loss: 0.0518 - val_acc: 0.9914

Epoch 3/100

- 144s - loss: 0.0261 - acc: 0.9943 - val_loss: 0.0554 - val_acc: 0.9905

Epoch 4/100

- 143s - loss: 0.0269 - acc: 0.9948 - val_loss: 0.0533 - val_acc: 0.9892

Epoch 00004: ReduceLROnPlateau reducing learning rate to 1.9999999494757503e-05.

Epoch 5/100

- 143s - loss: 0.0080 - acc: 0.9984 - val_loss: 0.0374 - val_acc: 0.9930

Epoch 6/100

- 143s - loss: 0.0030 - acc: 0.9992 - val_loss: 0.0491 - val_acc: 0.9936

Epoch 7/100

- 143s - loss: 0.0039 - acc: 0.9994 - val_loss: 0.0528 - val_acc: 0.9937

Epoch 00007: ReduceLROnPlateau reducing learning rate to 3.999999898951501e-06.

Epoch 8/100

- 143s - loss: 0.0015 - acc: 0.9997 - val_loss: 0.0550 - val_acc: 0.9934

Epoch 9/100

- 143s - loss: 7.9653e-04 - acc: 0.9999 - val_loss: 0.0622 - val_acc: 0.9932

Epoch 10/100

- 143s - loss: 6.3383e-04 - acc: 0.9999 - val_loss: 0.0677 - val_acc: 0.9933

Epoch 00010: ReduceLROnPlateau reducing learning rate to 7.999999979801942e-07.

Epoch 11/100

- 143s - loss: 4.9966e-04 - acc: 0.9999 - val_loss: 0.0688 - val_acc: 0.9934

Epoch 00011: early stopping

```
In [110]: test_loss, test_acc = model.evaluate(X_te, y_test_labels)
print('Test Accuracy:%4f, Test Loss:%4f.' %(test_acc,test_loss))
# Make Prediction
pred_result = model.predict(X_te)
y_pred=[]
for i in range(np.shape(y_test)[0]):
    num = np.where(pred_result[i]==max(pred_result[i]))
    y_pred.append(num[0][0])
y_pred = np.transpose(y_pred)
# calculate accuracy
```

```

num_test = len(y_test)
num_correct = np.sum(y_pred == y_test)
print('Got %d / %d correct' % (num_correct, num_test))
print('Accuracy = %f' % (np.mean(y_test == y_pred)))
from sklearn.metrics import classification_report, confusion_matrix
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred,
                           target_names=list(label_dict.values()),digits=3))

```

10000/10000 [=====] - 3s 310us/step

Test Accuracy:0.995300, Test Loss:0.049209.

Got 9953 / 10000 correct

Accuracy = 0.995300

```

[[ 977    0    0    0    0    0    2    1    0    0]
 [   0 1130    1    2    0    0    0    1    1    0]
 [   0    0 1029    1    0    0    0    2    0    0]
 [   0    0    0 1006    0    3    0    0    1    0]
 [   0    0    0    0  977    0    1    0    0    4]
 [   0    0    0    5    0  886    1    0    0    0]
 [   4    2    0    0    0    0  952    0    0    0]
 [   0    2    1    0    0    0    0 1024    0    1]
 [   0    0    0    1    0    0    0    0  972    1]
 [   0    0    0    0    4    1    0    2    2 1000]]

```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.996	0.997	0.996	980
1	0.996	0.996	0.996	1135
2	0.998	0.997	0.998	1032
3	0.991	0.996	0.994	1010
4	0.996	0.995	0.995	982
5	0.996	0.993	0.994	892
6	0.996	0.994	0.995	958
7	0.994	0.996	0.995	1028
8	0.996	0.998	0.997	974
9	0.994	0.991	0.993	1009

accuracy			0.995	10000
macro avg	0.995	0.995	0.995	10000
weighted avg	0.995	0.995	0.995	10000

8.4 Predict MNIST by Trained Imagenet Weight in VGG-16

```

In [195]: # Create base model of VGG16
from keras.layers import Input
input_tensor = Input(shape=(IMG_HEIGHT, IMG_WIDTH, IMG_DEPTH))
conv_base = VGG16(weights='imagenet',include_top=False,

```

```

        input_tensor = input_tensor)

    type(conv_base)

Out[195]: '\n#conv_base.summary()\nmodel = Sequential()\nfor layer in conv_base.layers:\n

In [0]: # Extracting features
        train_features = conv_base.predict(np.array(train_X), batch_size=BATCH_SIZE, verbose=1)
        test_features = conv_base.predict(np.array(X_te), batch_size=BATCH_SIZE, verbose=1)
        val_features = conv_base.predict(np.array(valid_X), batch_size=BATCH_SIZE, verbose=1)

48000/48000 [=====] - 18s 381us/step
10000/10000 [=====] - 4s 381us/step
12000/12000 [=====] - 5s 383us/step


In [0]: np.savez("train_features", train_features, train_label)
        np.savez("test_features", test_features, y_test_labels)
        np.savez("val_features", val_features, valid_label)

In [0]: # Current shape of features
        print(train_features.shape, "\n", test_features.shape, "\n", val_features.shape)

(48000, 1, 1, 512)
(10000, 1, 1, 512)
(12000, 1, 1, 512)


In [0]: # Flatten extracted features
        train_features_flat = np.reshape(train_features, (48000, 1*1*512))
        test_features_flat = np.reshape(test_features, (10000, 1*1*512))
        val_features_flat = np.reshape(val_features, (12000, 1*1*512))

In [0]: from keras import models
        from keras.models import Model
        from keras import layers
        from keras import optimizers
        from keras import callbacks
        from keras.layers.advanced_activations import LeakyReLU
        model_c = models.Sequential()
        model_c.add(layers.Dense(512, activation='relu', input_dim=(1*1*512)))
        model_c.add(layers.LeakyReLU(alpha=0.1))
        model_c.add(layers.Dense(10, activation='softmax'))

In [0]: NB_TRAIN_SAMPLES = train_features_flat.shape[0]
        NB_VALIDATION_SAMPLES = val_features_flat.shape[0]
        NB_EPOCHS = 100

        # Compile the model.
        model_c.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])

```

```
In [0]: # Incorporating reduced learning and early stopping for callback
```

```
reduce_learning = callbacks.ReduceLROnPlateau(  
    monitor='val_loss',  
    factor=0.2,  
    patience=2,  
    verbose=1,  
    mode='auto',  
    epsilon=0.0001,  
    cooldown=2,  
    min_lr=0)
```

```
early_stopping = callbacks.EarlyStopping(  
    monitor='val_loss',  
    min_delta=0.0003,  
    patience=8,  
    verbose=1,  
    mode='auto')
```

```
callbacks = [reduce_learning, early_stopping]
```

```
/usr/local/lib/python3.6/dist-packages/keras/callbacks.py:1065: UserWarning: `epsilon` argument  
warnings.warn("`epsilon` argument is deprecated and "
```

```
In [0]: # Train the the model
```

```
mt=model_c.fit(  
    train_features_flat,  
    train_label,  
    epochs=NB_EPOCHS,  
    validation_data=(val_features_flat, valid_label),  
    callbacks=callbacks  
)
```

Train on 48000 samples, validate on 12000 samples

Epoch 1/100

48000/48000 [=====] - 5s 112us/step - loss: 0.3400 - acc: 0.8962 - val.

Epoch 2/100

48000/48000 [=====] - 5s 97us/step - loss: 0.1631 - acc: 0.9472 - val.

Epoch 3/100

48000/48000 [=====] - 5s 97us/step - loss: 0.1389 - acc: 0.9547 - val.

Epoch 4/100

48000/48000 [=====] - 5s 97us/step - loss: 0.1198 - acc: 0.9608 - val.

Epoch 5/100

48000/48000 [=====] - 5s 95us/step - loss: 0.1087 - acc: 0.9643 - val.

Epoch 6/100

48000/48000 [=====] - 5s 96us/step - loss: 0.1028 - acc: 0.9663 - val.

Epoch 7/100

48000/48000 [=====] - 5s 97us/step - loss: 0.0941 - acc: 0.9685 - val.

```

Epoch 8/100
48000/48000 [=====] - 5s 97us/step - loss: 0.0897 - acc: 0.9692 - val.
Epoch 9/100
48000/48000 [=====] - 5s 97us/step - loss: 0.0842 - acc: 0.9719 - val.

Epoch 00009: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.
Epoch 10/100
48000/48000 [=====] - 5s 97us/step - loss: 0.0556 - acc: 0.9821 - val.
Epoch 11/100
48000/48000 [=====] - 5s 97us/step - loss: 0.0524 - acc: 0.9832 - val.
Epoch 12/100
48000/48000 [=====] - 5s 96us/step - loss: 0.0505 - acc: 0.9837 - val.
Epoch 13/100
48000/48000 [=====] - 5s 97us/step - loss: 0.0493 - acc: 0.9841 - val.
Epoch 14/100
48000/48000 [=====] - 5s 97us/step - loss: 0.0478 - acc: 0.9847 - val.
Epoch 15/100
48000/48000 [=====] - 5s 96us/step - loss: 0.0457 - acc: 0.9854 - val.
Epoch 16/100
48000/48000 [=====] - 5s 97us/step - loss: 0.0447 - acc: 0.9860 - val.
Epoch 17/100
48000/48000 [=====] - 5s 107us/step - loss: 0.0432 - acc: 0.9863 - val.
Epoch 18/100
48000/48000 [=====] - 5s 106us/step - loss: 0.0425 - acc: 0.9867 - val.
Epoch 19/100
48000/48000 [=====] - 5s 96us/step - loss: 0.0410 - acc: 0.9866 - val.
Epoch 20/100
48000/48000 [=====] - 5s 97us/step - loss: 0.0403 - acc: 0.9875 - val.
Epoch 21/100
48000/48000 [=====] - 5s 96us/step - loss: 0.0385 - acc: 0.9880 - val.

Epoch 00021: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.
Epoch 22/100
48000/48000 [=====] - 5s 97us/step - loss: 0.0333 - acc: 0.9901 - val.
Epoch 23/100
48000/48000 [=====] - 5s 96us/step - loss: 0.0326 - acc: 0.9907 - val.
Epoch 24/100
48000/48000 [=====] - 5s 96us/step - loss: 0.0322 - acc: 0.9908 - val.

Epoch 00024: ReduceLROnPlateau reducing learning rate to 8.000000525498762e-06.
Epoch 25/100
48000/48000 [=====] - 5s 95us/step - loss: 0.0310 - acc: 0.9915 - val.
Epoch 26/100
48000/48000 [=====] - 5s 95us/step - loss: 0.0309 - acc: 0.9915 - val.
Epoch 27/100
48000/48000 [=====] - 5s 96us/step - loss: 0.0308 - acc: 0.9916 - val.
Epoch 28/100
48000/48000 [=====] - 5s 95us/step - loss: 0.0307 - acc: 0.9915 - val.

```

```

Epoch 29/100
48000/48000 [=====] - 5s 96us/step - loss: 0.0307 - acc: 0.9918 - val.

Epoch 00029: ReduceLROnPlateau reducing learning rate to 1.6000001778593287e-06.
Epoch 30/100
48000/48000 [=====] - 5s 101us/step - loss: 0.0304 - acc: 0.9918 - val.
Epoch 31/100
48000/48000 [=====] - 5s 113us/step - loss: 0.0304 - acc: 0.9918 - val.
Epoch 32/100
48000/48000 [=====] - 5s 106us/step - loss: 0.0304 - acc: 0.9919 - val.

Epoch 00032: ReduceLROnPlateau reducing learning rate to 3.200000264769187e-07.
Epoch 33/100
48000/48000 [=====] - 5s 100us/step - loss: 0.0303 - acc: 0.9919 - val.
Epoch 34/100
48000/48000 [=====] - 5s 108us/step - loss: 0.0303 - acc: 0.9919 - val.
Epoch 35/100
48000/48000 [=====] - 5s 105us/step - loss: 0.0303 - acc: 0.9920 - val.

Epoch 00035: ReduceLROnPlateau reducing learning rate to 6.400000529538374e-08.
Epoch 36/100
48000/48000 [=====] - 5s 97us/step - loss: 0.0303 - acc: 0.9920 - val.
Epoch 37/100
48000/48000 [=====] - 5s 98us/step - loss: 0.0303 - acc: 0.9920 - val.
Epoch 00037: early stopping

```

```

In [0]: # Evaluate accuracy
        test_loss, test_acc = model_c.evaluate(test_features_flat, y_test_labels)
        print('Test Accuracy:%4f, Test Loss:%4f.' %(test_acc,test_loss))

10000/10000 [=====] - 0s 36us/step
Test Accuracy:0.978800, Test Loss:0.067221.

```

```

In [0]: # Make Prediction
        pred_result = model_c.predict(test_features_flat)
        y_pred=[]
        for i in range(np.shape(y_test)[0]):
            num = np.where(pred_result[i]==max(pred_result[i]))
            y_pred.append(num[0][0])
        y_pred = np.transpose(y_pred)

```

```

In [0]: # calculate accuracy
        num_test = len(y_test)
        num_correct = np.sum(y_pred == y_test)
        print('Got %d / %d correct' % (num_correct, num_test))
        print('Accuracy = %f' % (np.mean(y_test == y_pred)))
        from sklearn.metrics import classification_report, confusion_matrix

```

```

print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred,
                           target_names=list(label_dict.values()),digits=3))

```

Got 9788 / 10000 correct

Accuracy = 0.978800

```

[[ 969   0   1   0   0   4   5   1   0   0]
 [   0 1123   0   0   3   0   4   4   1   0]
 [   1   1 1008   4   2   4   4   5   2   1]
 [   0   0  10 977   0  16   0   3   2   2]
 [   0   1   1   0 967   1   2   3   4   3]
 [   2   0   7   9   1 864   0   2   6   1]
 [   7   0   4   0   0   2 944   0   1   0]
 [   0   5   4   2   8   0   0 1004   1   4]
 [   0   0   6   2   3   4   1   1 956   1]
 [   3   0   3   2   7   2   0   5  11 976]]

```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.987	0.989	0.988	980
1	0.994	0.989	0.992	1135
2	0.966	0.977	0.971	1032
3	0.981	0.967	0.974	1010
4	0.976	0.985	0.980	982
5	0.963	0.969	0.966	892
6	0.983	0.985	0.984	958
7	0.977	0.977	0.977	1028
8	0.972	0.982	0.977	974
9	0.988	0.967	0.977	1009

accuracy			0.979	10000
macro avg	0.979	0.979	0.979	10000
weighted avg	0.979	0.979	0.979	10000

8.5 Freeze All The Convolutional Layers Retrain VGG-16 Network On MNIST

```

In [0]: # Create base model of VGG16
        from keras.applications import VGG16;
        from keras.layers import Input
        input_tensor = Input(shape=(IMG_HEIGHT, IMG_WIDTH, IMG_DEPTH))
        freeze = VGG16(weights='imagenet',include_top=False,
                           input_tensor = input_tensor)
        freeze.summary()

```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

input_4 (InputLayer)	(None, 34, 34, 3)	0

block1_conv1 (Conv2D)	(None, 34, 34, 64)	1792

block1_conv2 (Conv2D)	(None, 34, 34, 64)	36928

block1_pool (MaxPooling2D)	(None, 17, 17, 64)	0

block2_conv1 (Conv2D)	(None, 17, 17, 128)	73856

block2_conv2 (Conv2D)	(None, 17, 17, 128)	147584

block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0

block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168

block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080

block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080

block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0

block4_conv1 (Conv2D)	(None, 4, 4, 512)	1180160

block4_conv2 (Conv2D)	(None, 4, 4, 512)	2359808

block4_conv3 (Conv2D)	(None, 4, 4, 512)	2359808

block4_pool (MaxPooling2D)	(None, 2, 2, 512)	0

block5_conv1 (Conv2D)	(None, 2, 2, 512)	2359808

block5_conv2 (Conv2D)	(None, 2, 2, 512)	2359808

block5_conv3 (Conv2D)	(None, 2, 2, 512)	2359808

block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

```

In [0]: # Freeze the layers which you don't want to train. Here I am freezing the first 5 layers
        for layer in freeze.layers:
            layer.trainable = False

Out[0]: '\n    for layer in freeze.layers[3:4]:\n        layer.trainable = False\nfor layer in freeze.layers[3:4]:\n    layer.trainable = False\n'

```



```

In [0]: # Extracting features
        train_features = freeze.predict(np.array(train_X), batch_size=BATCH_SIZE, verbose=1)
        test_features = freeze.predict(np.array(X_te), batch_size=BATCH_SIZE, verbose=1)
        val_features = freeze.predict(np.array(valid_X), batch_size=BATCH_SIZE, verbose=1)

48000/48000 [=====] - 18s 383us/step
10000/10000 [=====] - 4s 383us/step
12000/12000 [=====] - 5s 383us/step

In [0]: np.savez("train_features", train_features, train_label)
        np.savez("test_features", test_features, y_test_labels)
        np.savez("val_features", val_features, valid_label)

In [0]: # Current shape of features
        print(train_features.shape, "\n", test_features.shape, "\n", val_features.shape)

(48000, 1, 1, 512)
(10000, 1, 1, 512)
(12000, 1, 1, 512)

In [0]: # Flatten extracted features
        train_features_flat = np.reshape(train_features, (48000, 1*1*512))
        test_features_flat = np.reshape(test_features, (10000, 1*1*512))
        val_features_flat = np.reshape(val_features, (12000, 1*1*512))

In [0]: from keras import models
        from keras.models import Model
        from keras import layers
        from keras import optimizers
        from keras import callbacks
        from keras.layers.advanced_activations import LeakyReLU
        frz_con = models.Sequential()
        frz_con.add(layers.Dense(512, activation='relu', input_dim=(1*1*512)))
        frz_con.add(layers.LeakyReLU(alpha=0.1))
        frz_con.add(layers.Dense(512, activation='relu', input_dim=(1*1*512)))
        frz_con.add(layers.Dense(10, activation='softmax'))

In [0]: NB_TRAIN_SAMPLES = train_features_flat.shape[0]
        NB_VALIDATION_SAMPLES = val_features_flat.shape[0]
        NB_EPOCHS = 100

        # Compile the model.
        frz_con.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])

In [0]: # Incorporating reduced learning and early stopping for callback
        reduce_learning = callbacks.ReduceLROnPlateau(

```

```

        monitor='val_loss',
        factor=0.2,
        patience=2,
        verbose=1,
        mode='auto',
        epsilon=0.0001,
        cooldown=2,
        min_lr=0)

    eary_stopping = callbacks.EarlyStopping(
        monitor='val_loss',
        min_delta=0,
        patience=8,
        verbose=1,
        mode='auto')

    callbacks = [reduce_learning, eary_stopping]

/usr/local/lib/python3.6/dist-packages/keras/callbacks.py:1065: UserWarning: `epsilon` argument
  warnings.warn("`epsilon` argument is deprecated and ")

```

In [0]: # Train the the model

```

mt=frz_con.fit(
    train_features_flat,
    train_label,
    epochs=NB_EPOCHS,
    validation_data=(val_features_flat, valid_label),
    callbacks=callbacks
)

Train on 48000 samples, validate on 12000 samples
Epoch 1/100
48000/48000 [=====] - 6s 121us/step - loss: 0.3069 - acc: 0.8986 - va
Epoch 2/100
48000/48000 [=====] - 5s 112us/step - loss: 0.1606 - acc: 0.9480 - va
Epoch 3/100
48000/48000 [=====] - 6s 119us/step - loss: 0.1342 - acc: 0.9561 - va
Epoch 4/100
48000/48000 [=====] - 6s 121us/step - loss: 0.1205 - acc: 0.9593 - va
Epoch 5/100
48000/48000 [=====] - 5s 110us/step - loss: 0.1077 - acc: 0.9648 - va
Epoch 6/100
48000/48000 [=====] - 5s 111us/step - loss: 0.0987 - acc: 0.9674 - va

Epoch 00006: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.
Epoch 7/100
48000/48000 [=====] - 5s 111us/step - loss: 0.0571 - acc: 0.9815 - va

```

Epoch 8/100
 48000/48000 [=====] - 5s 112us/step - loss: 0.0503 - acc: 0.9838 - va
 Epoch 9/100
 48000/48000 [=====] - 5s 111us/step - loss: 0.0476 - acc: 0.9844 - va
 Epoch 10/100
 48000/48000 [=====] - 5s 111us/step - loss: 0.0447 - acc: 0.9853 - va
 Epoch 11/100
 48000/48000 [=====] - 5s 110us/step - loss: 0.0409 - acc: 0.9861 - va
 Epoch 12/100
 48000/48000 [=====] - 5s 111us/step - loss: 0.0387 - acc: 0.9874 - va

 Epoch 00012: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.
 Epoch 13/100
 48000/48000 [=====] - 5s 110us/step - loss: 0.0288 - acc: 0.9914 - va
 Epoch 14/100
 48000/48000 [=====] - 5s 111us/step - loss: 0.0271 - acc: 0.9917 - va
 Epoch 15/100
 48000/48000 [=====] - 5s 111us/step - loss: 0.0264 - acc: 0.9920 - va

 Epoch 00015: ReduceLROnPlateau reducing learning rate to 8.000000525498762e-06.
 Epoch 16/100
 48000/48000 [=====] - 5s 112us/step - loss: 0.0243 - acc: 0.9933 - va
 Epoch 17/100
 48000/48000 [=====] - 5s 111us/step - loss: 0.0239 - acc: 0.9933 - va
 Epoch 18/100
 48000/48000 [=====] - 6s 120us/step - loss: 0.0237 - acc: 0.9936 - va
 Epoch 19/100
 48000/48000 [=====] - 6s 120us/step - loss: 0.0236 - acc: 0.9934 - va

 Epoch 00019: ReduceLROnPlateau reducing learning rate to 1.6000001778593287e-06.
 Epoch 20/100
 48000/48000 [=====] - 5s 110us/step - loss: 0.0231 - acc: 0.9936 - va
 Epoch 21/100
 48000/48000 [=====] - 5s 111us/step - loss: 0.0230 - acc: 0.9936 - va
 Epoch 22/100
 48000/48000 [=====] - 5s 111us/step - loss: 0.0230 - acc: 0.9937 - va

 Epoch 00022: ReduceLROnPlateau reducing learning rate to 3.200000264769187e-07.
 Epoch 23/100
 48000/48000 [=====] - 5s 112us/step - loss: 0.0229 - acc: 0.9938 - va
 Epoch 24/100
 48000/48000 [=====] - 5s 110us/step - loss: 0.0229 - acc: 0.9938 - va
 Epoch 25/100
 48000/48000 [=====] - 5s 110us/step - loss: 0.0228 - acc: 0.9938 - va

 Epoch 00025: ReduceLROnPlateau reducing learning rate to 6.400000529538374e-08.
 Epoch 00025: early stopping

```
In [0]: # Evaluate accuracy
        test_loss, test_acc = frz_con.evaluate(test_features_flat, y_test_labels)
        print('Test Accuracy:%4f, Test Loss:%4f.' %(test_acc,test_loss))
```

10000/10000 [=====] - 0s 40us/step

Test Accuracy:0.977100, Test Loss:0.069289.

```
In [0]: # Make Prediction
        pred_result = frz_con.predict(test_features_flat)
        y_pred=[]
        for i in range(np.shape(y_test)[0]):
            num = np.where(pred_result[i]==max(pred_result[i]))
            y_pred.append(num[0][0])
        y_pred = np.transpose(y_pred)
```

```
In [0]: # calculate accuracy
        num_test = len(y_test)
        num_correct = np.sum(y_pred == y_test)
        print('Got %d / %d correct' % (num_correct, num_test))
        print('Accuracy = %f' % (np.mean(y_test == y_pred)))
        from sklearn.metrics import classification_report, confusion_matrix
        print(confusion_matrix(y_test, y_pred))
        print(classification_report(y_test, y_pred,
                                     target_names=list(label_dict.values()),digits=3))
```

Got 9771 / 10000 correct

Accuracy = 0.977100

```
[[ 967   0   1   0   0   3   5   1   2   1]
 [  0 1125   0   0   3   0   4   2   1   0]
 [  1   2 1000   8   2   6   3   6   3   1]
 [  0   0   9 976   0  18   0   4   2   1]
 [  0   1   0   0 968   1   3   2   3   4]
 [  2   0   4  13   1 861   3   2   5   1]
 [  7   1   2   0   1   3 941   0   2   1]
 [  0   3   4   3  10   0   0 1002   1   5]
 [  0   0   3   4   2   4   2   2 955   2]
 [  3   0   3   3   5   3   0   4  12 976]]
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.987	0.987	0.987	980
1	0.994	0.991	0.993	1135
2	0.975	0.969	0.972	1032
3	0.969	0.966	0.968	1010
4	0.976	0.986	0.981	982
5	0.958	0.965	0.961	892
6	0.979	0.982	0.981	958
7	0.978	0.975	0.976	1028
8	0.969	0.980	0.974	974

	9	0.984	0.967	0.976	1009
accuracy				0.977	10000
macro avg		0.977	0.977	0.977	10000
weighted avg		0.977	0.977	0.977	10000

8.6 Freeze All The Full Connected Layers Retrain VGG-16 Network On MNIST

```
In [0]: # Create base model of VGG16
        from keras.applications import VGG16;
        from keras.layers import Input
        input_tensor = Input(shape=(IMG_HEIGHT, IMG_WIDTH, IMG_DEPTH))
        freeze = VGG16(weights='imagenet',include_top=False,
                           input_tensor = input_tensor)
        freeze.summary()
```

Layer (type)	Output Shape	Param #
input_9 (InputLayer)	(None, 34, 34, 3)	0
block1_conv1 (Conv2D)	(None, 34, 34, 64)	1792
block1_conv2 (Conv2D)	(None, 34, 34, 64)	36928
block1_pool (MaxPooling2D)	(None, 17, 17, 64)	0
block2_conv1 (Conv2D)	(None, 17, 17, 128)	73856
block2_conv2 (Conv2D)	(None, 17, 17, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
block4_conv1 (Conv2D)	(None, 4, 4, 512)	1180160
block4_conv2 (Conv2D)	(None, 4, 4, 512)	2359808
block4_conv3 (Conv2D)	(None, 4, 4, 512)	2359808

```

-----
block4_pool (MaxPooling2D)      (None, 2, 2, 512)      0
-----
block5_conv1 (Conv2D)           (None, 2, 2, 512)      2359808
-----
block5_conv2 (Conv2D)           (None, 2, 2, 512)      2359808
-----
block5_conv3 (Conv2D)           (None, 2, 2, 512)      2359808
-----
block5_pool (MaxPooling2D)      (None, 1, 1, 512)      0
=====
Total params: 14,714,688
Trainable params: 14,714,688
Non-trainable params: 0
-----

```

```
In [0]: # Extracting features
```

```

train_features = freeze.predict(np.array(train_X), batch_size=BATCH_SIZE, verbose=1)
test_features = freeze.predict(np.array(X_te), batch_size=BATCH_SIZE, verbose=1)
val_features = freeze.predict(np.array(valid_X), batch_size=BATCH_SIZE, verbose=1)

```

```

48000/48000 [=====] - 19s 401us/step
10000/10000 [=====] - 4s 400us/step
12000/12000 [=====] - 5s 403us/step

```

```
In [0]: np.savez("train_features", train_features, train_label)
np.savez("test_features", test_features, y_test_labels)
np.savez("val_features", val_features, valid_label)

```

```
In [0]: # Current shape of features
```

```
print(train_features.shape, "\n", test_features.shape, "\n", val_features.shape)
```

```

(48000, 1, 1, 512)
(10000, 1, 1, 512)
(12000, 1, 1, 512)

```

```
In [0]: # Flatten extracted features
```

```

train_features_flat = np.reshape(train_features, (48000, 1*1*512))
test_features_flat = np.reshape(test_features, (10000, 1*1*512))
val_features_flat = np.reshape(val_features, (12000, 1*1*512))

```

```
In [0]: from keras import models
from keras.models import Model
from keras import layers
from keras import optimizers
from keras import callbacks

```

```

from keras.layers.advanced_activations import LeakyReLU
frz_fcn = models.Sequential()
frz_fcn.add(layers.Dense(512, activation='relu', input_dim=(1*1*512)))
frz_con.add(layers.Dense(512, activation='relu', input_dim=(1*1*512)))
frz_fcn.add(layers.Dense(10, activation='softmax'))

In [0]: # Freeze the layers which you don't want to train. Here I am freezing the first 5 layers
for layer in frz_fcn.layers[0:2]:
    layer.trainable = False

In [0]: NB_TRAIN_SAMPLES = train_features_flat.shape[0]
NB_VALIDATION_SAMPLES = val_features_flat.shape[0]
NB_EPOCHS = 100

# Compile the model.

frz_fcn.compile(optimizer='adam',
                loss='categorical_crossentropy',
                metrics=['accuracy'])

In [0]: # Incorporating reduced learning and early stopping for callback
reduce_learning = callbacks.ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.2,
    patience=2,
    verbose=1,
    mode='auto',
    epsilon=0.0001,
    cooldown=2,
    min_lr=0)

early_stopping = callbacks.EarlyStopping(
    monitor='val_loss',
    min_delta=0,
    patience=8,
    verbose=1,
    mode='auto')

callbacks = [reduce_learning, early_stopping]

/usr/local/lib/python3.6/dist-packages/keras/callbacks.py:1065: UserWarning: `epsilon` argument
warnings.warn("`epsilon` argument is deprecated and '

In [0]: # Train the the model
mt=frz_fcn.fit(
    train_features_flat,
    train_label,
    epochs=NB_EPOCHS,

```

```

        validation_data=(val_features_flat, valid_label),
        callbacks=callbacks
    )

Train on 48000 samples, validate on 12000 samples
Epoch 1/100
48000/48000 [=====] - 4s 73us/step - loss: 2.4916 - acc: 0.1330 - val.
Epoch 2/100
48000/48000 [=====] - 3s 55us/step - loss: 2.4916 - acc: 0.1330 - val.
Epoch 3/100
48000/48000 [=====] - 3s 55us/step - loss: 2.4916 - acc: 0.1330 - val.

Epoch 00003: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.
Epoch 4/100
48000/48000 [=====] - 3s 61us/step - loss: 2.4916 - acc: 0.1330 - val.
Epoch 5/100
48000/48000 [=====] - 3s 64us/step - loss: 2.4916 - acc: 0.1330 - val.
Epoch 6/100
48000/48000 [=====] - 3s 54us/step - loss: 2.4916 - acc: 0.1330 - val.

Epoch 00006: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.
Epoch 7/100
48000/48000 [=====] - 3s 54us/step - loss: 2.4916 - acc: 0.1330 - val.
Epoch 8/100
48000/48000 [=====] - 3s 55us/step - loss: 2.4916 - acc: 0.1330 - val.
Epoch 9/100
48000/48000 [=====] - 3s 54us/step - loss: 2.4916 - acc: 0.1330 - val.

Epoch 00009: ReduceLROnPlateau reducing learning rate to 8.000000525498762e-06.
Epoch 00009: early stopping

In [0]: # Evaluate accuracy
        test_loss, test_acc = frz_fcn.evaluate(test_features_flat, y_test_labels)
        print('Test Accuracy:%4f, Test Loss:%4f.' %(test_acc,test_loss))

10000/10000 [=====] - 0s 43us/step
Test Accuracy:0.129600, Test Loss:2.491812.

In [0]: # Make Prediction
        pred_result = frz_fcn.predict(test_features_flat)
        y_pred=[]
        for i in range(np.shape(y_test)[0]):
            num = np.where(pred_result[i]==max(pred_result[i]))
            y_pred.append(num[0][0])
        y_pred = np.transpose(y_pred)

In [0]: # calculate accuracy
        num_test = len(y_test)

```



```

num_correct = np.sum(y_pred == y_test)
print('Got %d / %d correct' % (num_correct, num_test))
print('Accuracy = %f' % (np.mean(y_test == y_pred)))
from sklearn.metrics import classification_report, confusion_matrix
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred,
                           target_names=list(label_dict.values()),digits=3))

```

Got 1296 / 10000 correct

Accuracy = 0.129600

```

[[ 0  0  0  0  0  0  0  0  0 980  0]
 [ 0  0  0  4 821 266  0  2  34  8]
 [ 0  0  0  0  5  0  0 10 1015  2]
 [ 0  0  0  0  4  0  0  0 1005  1]
 [ 0  0  0  5 315 27  0 12  612 11]
 [ 0  0  0  0  8  0  0  2  881  1]
 [ 0  0  0  0  8  1  0  3  939  7]
 [ 0  0  0  0 217 48  0  8  702 53]
 [ 0  0  0  0  2  0  0  0  972  0]
 [ 0  0  0  0 15  5  0  0  988 11]]

```

	precision	recall	f1-score	support
0	0.000	0.000	0.000	980
1	0.000	0.000	0.000	1135
2	0.000	0.000	0.000	1032
3	0.000	0.000	0.000	1010
4	0.226	0.321	0.265	982
5	0.000	0.000	0.000	892
6	0.000	0.000	0.000	958
7	0.216	0.008	0.015	1028
8	0.120	0.998	0.214	974
9	0.012	0.001	0.002	1009
accuracy				0.130 10000
macro avg				0.057 0.133 0.050 10000
weighted avg				0.057 0.130 0.049 10000

```

/usr/local/lib/python3.6/dist-packages/sklearn/metrics/classification.py:1437: UndefinedMetricWarning: Precision is undefined because no sample was correctly classified
'precision', 'predicted', average, warn_for)

```